Why might a process be placed on the ready queue?

What is 'wait time'? Total wait time, or the first waiting before it is scheduled the first time?

Write a formula for the wait time based on arrival time, execution time(=duration) and completion time

Scheduling. Some terms... How shall I compare thee?

"Turnaround time"

"Waiting time"

"Response time"

"Throughput"

"Latency"

"Starvation?

Good for Batch? Good use of CPU/IO resources?

Good for Interactive?

Good for real-time systems?

FCFS (aka.....)

SJF vs Pre-emptive SJF

RR

Priority-scheduling

Choosing an appropriate time-quantum. What does scheduler does Linux use?

Which schedulers can suffer from starvation?

Which schedulers are appropriate for batch jobs?

Which schedulers are appropriate for interactive jobs?

What scheduler does Linux use?

Determine the scheduling sequence and calculate the average wait time of the following schedulers
Tie-break #1: Schedule the earliest arriving job. Tie break #2: P4 is placed on ready queue first

Round robin (quanta = 10ms)

Process	Arrival Time(ms)	Burst Time(ms)	Wait Time (ms)
P1	0	30	
P2	0	20	
P3	0	20	
P4	10	10	

0..10	.. 20	..30	..40	..50	.. 60	.. 70	..80

Shortest Job First

Process	Arrival Time(ms)	Burst Time(ms)	Wait Time (ms)
P1	0	30	
P2	0	20	
P3	0	20	
P4	10	10	

0..10	20	30	40	50	60	70	80

First Come First Served (assume arrive in order P1,P2,P3)

Process	Arrival Time(ms)	Burst Time(ms)	Wait Time (ms)
P1	0	30	
P2	0	20	
P3	0	20	
P4	10	10	

0..10	..20	30	40	50	60	70	80

Pre-emptive Shortest Job First (assume interrupted jobs are placed at the front of the queue)

Process	Arrival T	Burst T	Wait T
P1	0	30	
P2	0	20	
P3	0	20	
P4	10	10	

0..10	..20	30	40	50	60	70	80

Pre-emptive Priority (higher value = higher priority)

Process	Arrival	Burst	Priority	Wait
P1	0	30	2	
P2	0	20	4	
P3	0	20	1	
P4	10	10	3	

0..10	..20	30	40	50	60	70	80

What is the **Convoy Effect** (poor I/O parallelism)?